

Disciplined Nonlinear Programming

Daniel Cederberg

Based on joint work with William Zhang, Parth Nobel, and Stephen Boyd

March 2026 ¹

¹Versions of this talk have been given at Stanford, KTH, and Amazon.

Outline

Convex optimization

Nonlinear programming

Disciplined nonlinear programming

Examples

Conclusions

Convex optimization

$$\begin{array}{ll} \text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq 0, \quad i = 1, \dots, m \\ & Ax = b \end{array}$$

- ▶ variable $x \in \mathbf{R}^n$
- ▶ equality constraints are linear and f_0, \dots, f_m are convex
- ▶ why care about convexity?

Convex optimization

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 0, \quad i = 1, \dots, m \\ & && Ax = b \end{aligned}$$

- ▶ variable $x \in \mathbf{R}^n$
- ▶ equality constraints are linear and f_0, \dots, f_m are convex
- ▶ why care about convexity?
 - **classic answer**: local optimum = global optimum
 - **modern answer**: very actionable because of reliable software

Disciplined convex programming (DCP)

- ▶ **motivation:** difficult to verify convexity of an arbitrarily constructed problem
- ▶ a grammar for specifying convex optimization problems [Grant et al., 2006]
- ▶ convex and concave functions are constructed using simple composition rules
- ▶ a problem that follows the DCP rules is convex-by-construction

Disciplined convex programming (DCP)

- ▶ **motivation:** difficult to verify convexity of an arbitrarily constructed problem
- ▶ a grammar for specifying convex optimization problems [Grant et al., 2006]
- ▶ convex and concave functions are constructed using simple composition rules
- ▶ a problem that follows the DCP rules is convex-by-construction
- ▶ and can be solved globally and efficiently (in theory and in practice)

Modeling languages implementing DCP

- ▶ Yalmip & CVX (MATLAB), CVXPY (Python), CVXR (R), Convex.jl (Julia)
- ▶ classic commercial software like AMPL and GAMS do not implement DCP

Modeling languages implementing DCP

- ▶ Yalmip & CVX (MATLAB), CVXPY (Python), CVXR (R), Convex.jl (Julia)
- ▶ classic commercial software like AMPL and GAMS do not implement DCP

- ▶ two main tasks of DCP-based modeling languages:
 - **analysis**: verify that a high level problem description follows the DCP rules
 - **synthesis**: build an equivalent problem on standard conic form:

$$\begin{array}{ll} \text{minimize} & c^T x \\ \text{subject to} & Ax + s = b \\ & s \in K \end{array}$$

variables x and s , and problem data A , b , c and convex cone K

Modeling languages implementing DCP

- ▶ Yalmip & CVX (MATLAB), CVXPY (Python), CVXR (R), Convex.jl (Julia)
- ▶ classic commercial software like AMPL and GAMS do not implement DCP

- ▶ two main tasks of DCP-based modeling languages:
 - **analysis**: verify that a high level problem description follows the DCP rules
 - **synthesis**: build an equivalent problem on standard conic form:

$$\begin{array}{ll} \text{minimize} & c^T x \\ \text{subject to} & Ax + s = b \\ & s \in K \end{array}$$

variables x and s , and problem data A , b , c and convex cone K

- ▶ no communication between modeling language and solver beyond A, b, c, K

Disciplined convex programming in practice: An example


$$\text{minimize } -\log \det X + \mathbf{Tr}(SX) + \lambda \|X\|_1$$

- ▶ variable $X \in \mathbf{S}^n$ represents an inverse covariance matrix of a Gaussian vector
- ▶ the problem data is the sample covariance matrix $S \in \mathbf{S}_+^n$
- ▶ called *sparse inverse covariance estimation* (applications in graphical modeling)

Disciplined convex programming in practice: An example

$$\text{minimize } -\log \det X + \text{Tr}(SX) + \lambda \|X\|_1$$

- ▶ variable $X \in \mathbf{S}^n$ represents an inverse covariance matrix of a Gaussian vector
- ▶ the problem data is the sample covariance matrix $S \in \mathbf{S}_+^n$
- ▶ called *sparse inverse covariance estimation* (applications in graphical modeling)



Disciplined convex programming in practice: An example

$$\text{minimize } -\log \det X + \text{Tr}(SX) + \lambda \|X\|_1$$

What the user sees:

```
1 import cvxpy as cp
2 S, lambda = ...
3 X = cp.Variable((n, n), symmetric=True)
4 obj = -cp.logdet(X) + cp.trace(S @ X) +
      lambda * cp.norm1(X)
5 prob = cp.Problem(cp.Minimize(obj))
6 prob.solve()
```

What the solver sees:

$$\begin{aligned} \text{minimize} \quad & c^T x \\ \text{subject to} \quad & Ax + s = b \\ & s \in K \end{aligned}$$

where A, b, c, K and x are defined on the next slide impossible to derive by hand

What do to when your problem is not convex?

- ▶ sometimes possible to reformulate as convex using *e.g.*, a change of variables
 - maximum likelihood estimation for exponential families
 - feedback control design
 - geometric programming
- ▶ implement simple heuristics based on convex optimization
 - sequential convex programming
 - convex-concave procedure
 - alternating convex optimization
- ▶ use a general-purpose nonlinear programming solver

Outline

Convex optimization

Nonlinear programming

Disciplined nonlinear programming

Examples

Conclusions


Nonlinear programming (NLP)

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & c(x) = 0 \\ & \ell \leq x \leq u \end{array}$$

- ▶ $f : \mathbf{R}^n \rightarrow \mathbf{R}$, $c : \mathbf{R}^n \rightarrow \mathbf{R}^p$ differentiable, not necessarily convex
- ▶ we seek a **local solution**


Misconception 1 about NLP

NLP solvers always compute global minimizers




Misconception 1 about NLP

NLP solvers always compute global minimizers



Misconception 1 about NLP


NLP solvers always compute global minimizers



- ▶ Ipopt [Wächter and Biegler, 2006] converges to a local (but not global) minimizer


Misconception 2 about NLP

NLP solvers always converge to the local minimizer closest to the initial point




Misconception 2 about NLP

NLP solvers always converge to the local minimizer closest to the initial point



Misconception 2 about NLP


NLP solvers always converge to the local minimizer closest to the initial point



- ▶ Ipopt converges to a local minimizer, but not the one closest to the initial point


Misconception 3 about NLP

NLP solvers always converge to a local minimizer




Misconception 3 about NLP

NLP solvers always converge to a local minimizer



Misconception 3 about NLP

NLP solvers always converge to a local minimizer



- ▶ Ipopt converges to a saddle point when minimizing $f(x) = x^3$ initialized at $x_0 = 1$

Misconception 4 about NLP


NLP solvers always converge to a 'good' feasible point (KKT point)

Misconception 4 about NLP

$$\begin{aligned} &\text{minimize} && f(x) = x^4 - 2x^2 + 0.5x + 1 \\ &\text{subject to} && f(x) \leq 0 \end{aligned}$$

(Also see Wächter & Bowly, Gurobi, 2026)

NLP solvers always converge to a 'good' feasible point (KKT point)




Misconception 4 about NLP

$$\begin{aligned} &\text{minimize} && f(x) = x^4 - 2x^2 + 0.5x + 1 \\ &\text{subject to} && f(x) \leq 0 \end{aligned}$$

(Also see Wächter & Bowly, Gurobi, 2026)

NLP solvers always converge to a 'good' feasible point (KKT point)




Misconception 4 about NLP

$$\begin{aligned} &\text{minimize} && f(x) = x^4 - 2x^2 + 0.5x + 1 \\ &\text{subject to} && f(x) \leq 0 \end{aligned}$$

(Also see Wächter & Bowly, Gurobi, 2026)

NLP solvers always converge to a 'good' feasible point (KKT point)




- ▶ Ipopt converges to an infeasible point when initialized at $x_0 = 0.8$

Misconception 4 about NLP

$$\begin{aligned} &\text{minimize} && f(x) = x^4 - 2x^2 + 0.5x + 1 \\ &\text{subject to} && f(x) \leq 0 \end{aligned}$$

(Also see Wächter & Bowly, Gurobi, 2026)

NLP solvers always converge to a 'good' feasible point (KKT point)




Misconception 4 about NLP

$$\begin{aligned} &\text{minimize} && f(x) = x^4 - 2x^2 + 0.5x + 1 \\ &\text{subject to} && f(x) \leq 0 \end{aligned}$$

(Also see Wächter & Bowly, Gurobi, 2026)

NLP solvers always converge to a 'good' feasible point (KKT point)



- ▶ Ipopt converges to global minimizer when initialized at $x_0 = 1.2$

Summary of misconceptions

NLP solvers are **not** guaranteed to

- ▶ compute global minimizers
- ▶ compute the local minimizer closest to the initial point
- ▶ compute local minimizers at all (might converge to saddle points)
- ▶ compute a feasible point at all (might converge to an infeasible point)

Summary of misconceptions

NLP solvers are **not** guaranteed to

- ▶ compute global minimizers
- ▶ compute the local minimizer closest to the initial point
- ▶ compute local minimizers at all (might converge to saddle points)
- ▶ compute a feasible point at all (might converge to an infeasible point)

but in practice, NLP solvers often work well (perhaps with decent initialization)

Constraint qualifications (CQs)

- ▶ NLP solvers attempt to compute **KKT points**
- ▶ KKT conditions rely on a local linear approximation of the feasible set
 - KKT is necessary for local optimality only when this approximation is “good enough”
 - (formally: the tangent cone equals the linearized tangent cone)
- ▶ **constraint qualifications** are conditions that guarantee this
 - most common CQ: LICQ (gradients of active constraints are linearly independent)

Constraint qualifications (CQs)

- ▶ NLP solvers attempt to compute **KKT points**
- ▶ KKT conditions rely on a local linear approximation of the feasible set
 - KKT is necessary for local optimality only when this approximation is “good enough”
 - (formally: the tangent cone equals the linearized tangent cone)
- ▶ **constraint qualifications** are conditions that guarantee this
 - most common CQ: LICQ (gradients of active constraints are linearly independent)
 - we found LICQ to be important in practice

Modeling languages for NLP

- ▶ YALMIP (MATLAB), JuMP (Julia), Pyomo (Python), and CasADi (C++)
- ▶ commercial software like AMPL, GAMS, AIMMS

Modeling languages for NLP

- ▶ YALMIP (MATLAB), JuMP (Julia), Pyomo (Python), and CasADi (C++)
- ▶ commercial software like AMPL, GAMS, AIMMS
- ▶ main task of NLP modeling languages:
 - provide solvers with derivatives
- ▶ communication between modeling language and solver in every iteration

When many existing NLP modeling languages do a bad job

$$\text{minimize } \|Ax - b\|_2^2 + \lambda \|x\|_1$$

- ▶ variable x , regularization parameter $\lambda > 0$
- ▶ convex problem (lasso) so should be solved gracefully
- ▶ the solution occurs at a point of nondifferentiability
- ▶ objective is treated as a black box and “derivatives” are supplied to the solver
- ▶ ignoring nondifferentiability might cause solvers to struggle

When existing NLP modeling languages do a bad job

output from Ipopt when interfacing it using Casadi:

```
This is Ipopt version 3.14.11, running with linear solver MUMPS 5.4.1.
```

```
...
```

```
iter    objective    inf_pr  inf_du lg(mu)  ||d||  lg(rg) alpha_du alpha_pr ls
  10  1.1126107e+01  0.00e+00  1.71e+00  -1.7  1.29e+08  -7.8  1.00e+00  1.32e-23f 77
  11  1.3152054e+10  0.00e+00  1.71e+00  -1.7  4.09e+08  -8.3  1.00e+00  1.00e+00w  1
  12  3.5423855e+10  0.00e+00  1.66e+00  -1.7  1.13e+09  -8.8  1.00e+00  1.00e+00w  1
  13  1.4112259e+11  0.00e+00  1.58e+00  -1.7  3.39e+09  -9.2  1.00e+00  1.00e+00w  1
```

```
Cannot call restoration phase at point that is almost feasible (violation 0.000000e+00).  
Abort in line search due to no other fall back.
```

```
Number of Iterations....: 13
```

```
EXIT: Error in step computation!
```

Outline

Convex optimization

Nonlinear programming

Disciplined nonlinear programming

Examples

Conclusions

Disciplined nonlinear programming (DNLP)

- ▶ a new grammar for specifying nonlinear programs
- ▶ smooth functions can be freely mixed with nonsmooth functions
- ▶ rules governing how nonsmooth functions can appear
- ▶ smooth = twice continuously differentiable in the interior of their domain

Disciplined nonlinear programming (DNLP)

- ▶ a new grammar for specifying nonlinear programs
- ▶ smooth functions can be freely mixed with nonsmooth functions
- ▶ rules governing how nonsmooth functions can appear
- ▶ smooth = twice continuously differentiable in the interior of their domain

a DNLP-compliant problem can be transformed to a form that

- ▶ only involves smooth functions
- ▶ does not introduce LICQ violations in the transformation

The DNLN grammar

in DNLN, problems are constructed using **atoms** with known attributes

- ▶ atoms are divided into three categories:
 - **smooth**
 - **nonsmooth convex**
 - **nonsmooth concave**
- ▶ DNLN enforces simple rules on how atoms can be composed

The DNLP grammar

in DNLP, problems are constructed using **atoms** with known attributes

- ▶ atoms are divided into three categories:
 - **smooth**
 - **nonsmooth convex**
 - **nonsmooth concave**
- ▶ DNLP enforces simple rules on how atoms can be composed
- ▶ the DNLP rules are very simple if you know the DCP rules:

a problem is DNLP if it is DCP when smooth atoms are linearized

Some new DNLN atoms

Atom	Definition
multiply	$\phi(x, y) = xy$
quad_form	$\phi(x) = x^T Qx$ where $Q \in \mathbf{S}^n$
sin	$\phi(x) = \sin x$
sigmoid	$\phi(x) = \frac{1}{1+e^{-x}}$
normal_cdf	$\phi(x) = (1/\sqrt{2\pi}) \int_{-\infty}^x e^{-t^2/2} dt$

- ▶ these atoms are not convex or concave, but they are smooth

Dealing with nonsmooth atoms

- ▶ a DNLP problem is equivalent to a problem where any
 - nonsmooth convex atom ϕ appears in a constraint of the form $\phi(x) \leq t$
 - nonsmooth concave atom ϕ appears in a constraint of the form $\phi(x) \geq t$

Dealing with nonsmooth atoms


- ▶ a DNLP problem is equivalent to a problem where any
 - nonsmooth convex atom ϕ appears in a constraint of the form $\phi(x) \leq t$
 - nonsmooth concave atom ϕ appears in a constraint of the form $\phi(x) \geq t$
- ▶ these constraints can be formulated using smooth functions [Grant and Boyd, 2008]
 - the reformulations do not introduce LICQ violations
- ▶ **example:** $|x| \leq t$ can be reformulated as $-t \leq x \leq t$

Dealing with nonsmooth atoms

- ▶ a DNLP problem is equivalent to a problem where any
 - nonsmooth convex atom ϕ appears in a constraint of the form $\phi(x) \leq t$
 - nonsmooth concave atom ϕ appears in a constraint of the form $\phi(x) \geq t$
- ▶ these constraints can be formulated using smooth functions [Grant and Boyd, 2008]
 - the reformulations do not introduce LICQ violations
- ▶ **example:** $|x| \leq t$ can be reformulated as $-t \leq x \leq t$
- ▶ other nonsmooth atoms, such as the Huber atom

$$\phi(x) = \begin{cases} x^2 & |x| \leq M \\ 2M|x| - M^2 & |x| > M \end{cases}$$

have more complex reformulations



Outline

Convex optimization

Nonlinear programming

Disciplined nonlinear programming

Examples

Conclusions

A DNLP extension of CVXPY

- ▶ DNLP is implemented in CVXPY version 1.9
- ▶ supported solvers:
 - IPOPT (open-source, interior-point method)
 - Uno [Vanaret and Leyffer, 2025] (open-source, IPM and SQP)
 - KNITRO (commercial, many algorithms)
 - COPT (commercial, interior-point method)
- ▶ computes sparse Jacobians and Hessians via C backend


Example: Circle packing

$$\begin{aligned} & \text{minimize} && \max_{i=1,\dots,n} (\|c_i\|_\infty + r_i) \\ & \text{subject to} && \|c_i - c_j\|_2^2 \geq (r_i + r_j)^2, \quad 1 \leq i < j \leq n, \end{aligned}$$

- ▶ goal is to pack n circles with given radii r_i into smallest square
- ▶ variables are the circle centers $c_i \in \mathbf{R}^2$
- ▶ becomes DCP when smooth atoms are linearized, so DNLP-compliant



Example: Circle packing

- ▶ result for $n = 10$ circles with random radii, with random initial point
- ▶ the fraction of the square covered by the circles is 0.72



Example: Circle packing

- ▶ result for same problem instance over 500 random initializations using `best_of=500`
- ▶ best coverage found is 0.77



Example: Sparse signal recovery

$$\begin{array}{ll} \text{minimize} & \sum_{i=1}^n \sqrt{|x_i|} \\ \text{subject to} & Ax = y, \end{array}$$

- ▶ goal is to recover a sparse signal $x_0 \in \mathbf{R}^n$ from m given measurements $y = Ax_0$
- ▶ variable is the signal $x \in \mathbf{R}^n$
- ▶ the objective is nonconvex but promotes sparsity better than the ℓ_1 norm
- ▶ becomes DCP when smooth atoms are linearized, so DNLP-compliant


Example: Sparse signal recovery

$$\begin{aligned} & \text{minimize} && \sum_{i=1}^n \sqrt{|x_i|} \\ & \text{subject to} && Ax = y, \end{aligned}$$

- ▶ goal is to recover a sparse signal $x_0 \in \mathbf{R}^n$ from m given measurements $y = Ax_0$
- ▶ variable is the signal $x \in \mathbf{R}^n$
- ▶ the objective is nonconvex but promotes sparsity better than the ℓ_1 norm
- ▶ becomes DCP when smooth atoms are linearized, so DNLP-compliant

```
1 x = Variable(n)
2 cost, constr = sum(sqrt(abs(x))), [A @ x == y]
3 prob = Problem(Minimize(cost), constr)
4 prob.solve(nlp=True, solver='knitro_ipm')
```

Example: Sparse signal recovery



- probability of successful recovery for signal dimension $n = 100$, nonconvex approach (left) and convex ℓ_1 approach (right)

Many more examples in the paper

- ▶ planning and geometry
- ▶ signal processing
- ▶ finance
- ▶ energy systems
- ▶ statistics

Outline

Convex optimization

Nonlinear programming

Disciplined nonlinear programming

Examples

Conclusions

Conclusions

- ▶ DNLP is a new grammar for specifying NLPs
 - no rules if all atoms are smooth
 - mitigates solver failures caused by nonsmoothness
- ▶ the DNLP-extension of CVXPY makes it easy to use NLP solvers
 - enables high level, human readable problem descriptions
 - automatically generates smooth reformulations of nonsmooth atoms
 - ideal for teaching and research